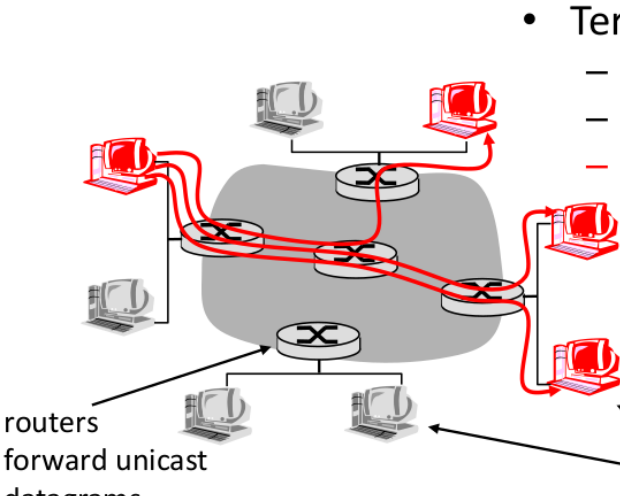


# Multicast



- Terminology
  - Unicast: one-to-one
  - Broadcast: one-to-all
  - **Multicast: one-to-many**

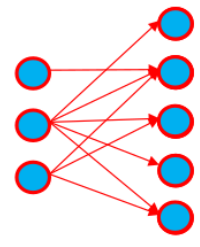
[Simulcast]

- Portmanteau of simultaneous broadcast
- Broadcasting of programs or events across more than one medium
- Also used in context of unicast (adaptive streaming)

multicast receiver (red)

not a multicast receiver

routers forward unicast datagrams

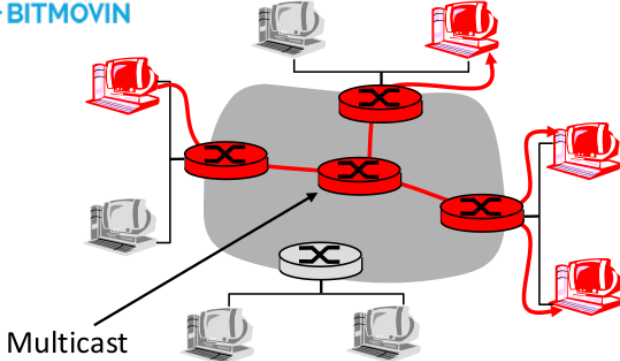


Any-Cast; Server → Klienten: z.B.: Klienten melden sich zu Stream an.

Arten von Multicast:

► BITMOVIN

ALPEN-ADRIA  
UNIVERSITÄT  
KLAUFERTECHNIK




Multicast routers (red) duplicate and forward multicast datagrams

- Application-layer multicast
  - End systems involved in multicast copy and forward unicast datagrams

## • Network multicast


- Routers actively participate in multicast, making copies of packets forwarding as needed



Network-Multicast: Auf IP-Schicht :: Im größeren Kontext nicht definiert.

Es gibt vordefinierte Multicast-Adressen:

# Internet Multicast Service Model



- Multicast group concept: use of indirection
  - Hosts can join a **multicast group**
  - Datagrams sent to the group are **forwarded** by the routers to the **members**
  - Group management is **not trivial**
  - Mainly used **within ISPs** (not across multiple ISPs), e.g., for IPTV

Obrige Klienten haben sich zu der Multicast-Gruppe angemeldet & Router weiß wohins geroutet wird.

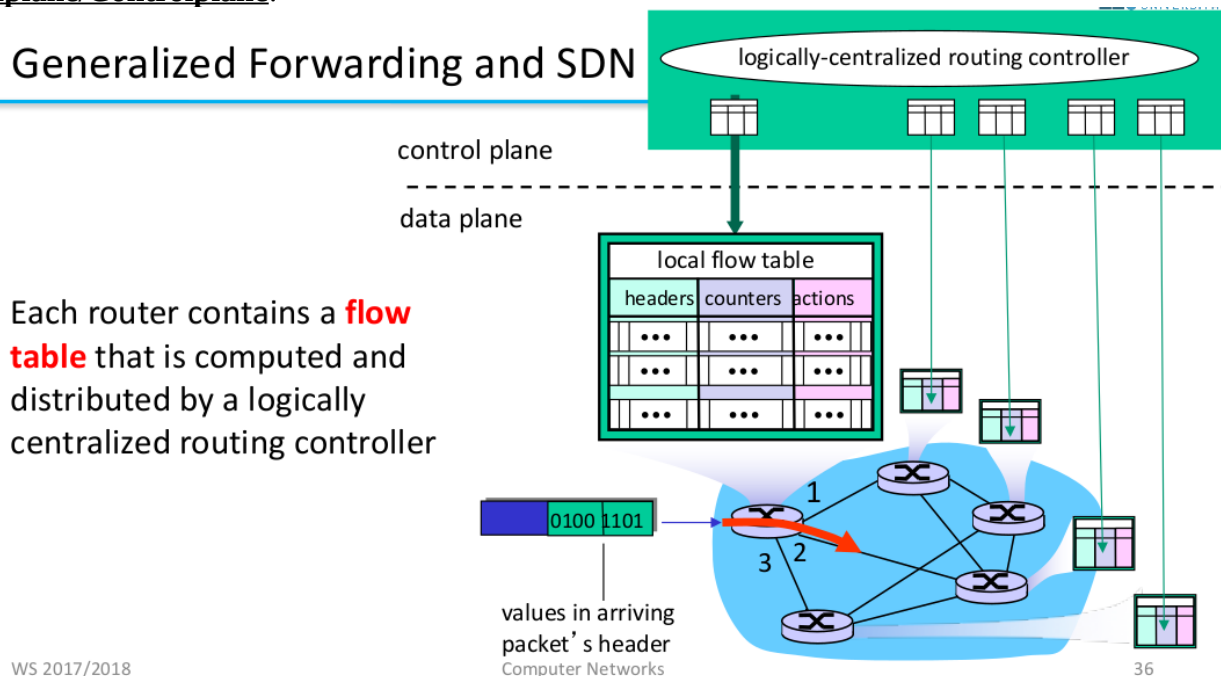
Classful IP Addressing:

D :: [1110] Multicast Address | ↔ 224.0.0.0 to 239.255.255.255

z.B.: Asfinag, Autobahnkameraüberwachung ist multicasted via Session Encrypted Protocol ::  
Kaum Sicherheit

**Dataplane/Controlplane:**

## Generalized Forwarding and SDN




Each router contains a **flow table** that is computed and distributed by a logically centralized routing controller

Local Flow Table :: Forwarding-Tabelle predefinierter Actions @Router

# OpenFlow Data Plane Abstraction

- **Flow**: defined by **header fields**
- **Generalized forwarding**: simple packet-handling rules
  - **Pattern**: match values in packet header fields
  - **Actions**: for matched packet: drop, forward, modify, matched packet or send matched packet to controller
  - **Priority**: disambiguate overlapping patterns
  - **Counters**: #bytes and #packets



\* : wildcard

Flow table in a router (computed and distributed by controller) define router's **match + action rules**

1. src=1.2.\*.\*, dest=3.4.5.\* → drop
2. src = \*.\*.\*.\*, dest=3.4.\*.\* → forward(2)
3. src=10.1.2.3, dest=.\*.\*.\* → send to controller

WS 2017/2018

Computer Networks

37

Flow: Definiert wohin was geht

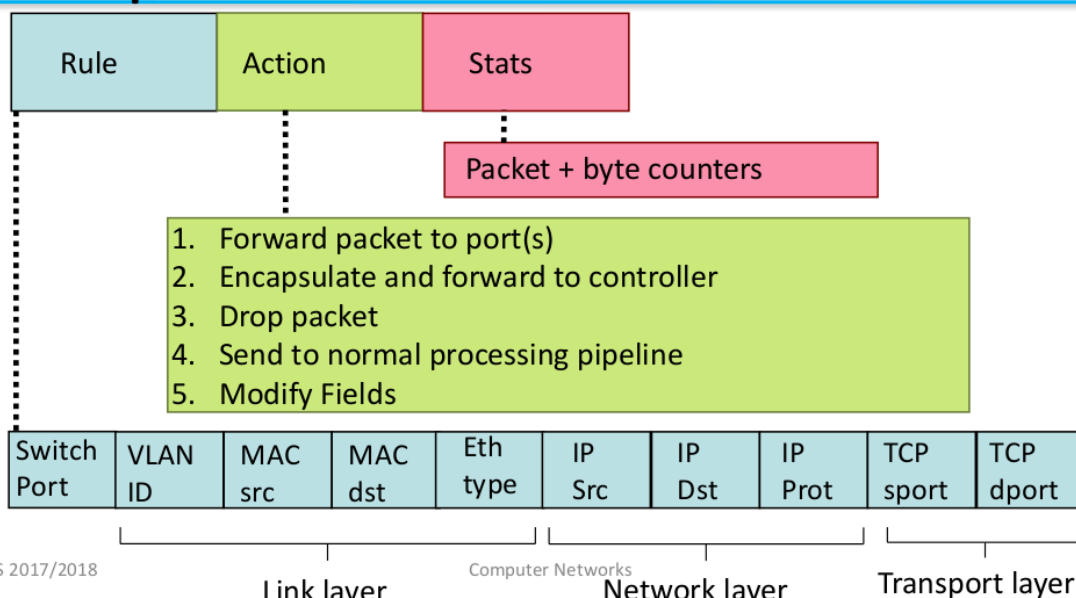
Gen. Forwarding: Predefiniert im Router

- Pattern matches Header
- Actions: .. taken
- Priority: Mehrere pattern matchen auf 1 Action: Prioritätsverweisung
- Counters: #Bytes #Packets

'Match + Action Rules' – Table above :: beliebig Programmierbar

Bsp.: 1 SDN-Framework:

## OpenFlow: Flow Table Entries



WS 2017/2018

Computer Networks

38

Frei

| Switch Port | MAC src | MAC dst | Eth type | VLAN ID | IP Src | IP Dst   | IP Prot | TCP sport | TCP dport | Action |
|-------------|---------|---------|----------|---------|--------|----------|---------|-----------|-----------|--------|
| *           | *       | *       | *        | *       | *      | 51.6.0.8 | *       | *         | *         | port6  |

*IP datagrams destined to IP address 51.6.0.8 should be forwarded to router output port 6*

### Firewall:

| Switch Port | MAC src | MAC dst | Eth type | VLAN ID | IP Src | IP Dst | IP Prot | TCP sport | TCP dport | Action |
|-------------|---------|---------|----------|---------|--------|--------|---------|-----------|-----------|--------|
| *           | *       | *       | *        | *       | *      | *      | *       | *         | 22        | drop   |

*do not forward (block) all datagrams destined to TCP port 22*

| Switch Port | MAC src | MAC dst | Eth type | VLAN ID | IP Src      | IP Dst | IP Prot | TCP sport | TCP dport | Action |
|-------------|---------|---------|----------|---------|-------------|--------|---------|-----------|-----------|--------|
| *           | *       | *       | *        | *       | 128.119.1.1 | *      | *       | *         | *         | drop   |

*do not forward (block) all datagrams sent by host 128.119.1.1*

WS 2017/2018

Computer Networks

39

## Examples

### Destination-based layer 2 (switch) forwarding:

| Switch Port | MAC src           | MAC dst | Eth type | VLAN ID | IP Src | IP Dst | IP Prot | TCP sport | TCP dport | Action |
|-------------|-------------------|---------|----------|---------|--------|--------|---------|-----------|-----------|--------|
| *           | 22:A7:23:11:E1:02 | *       | *        | *       | *      | *      | *       | *         | *         | port3  |

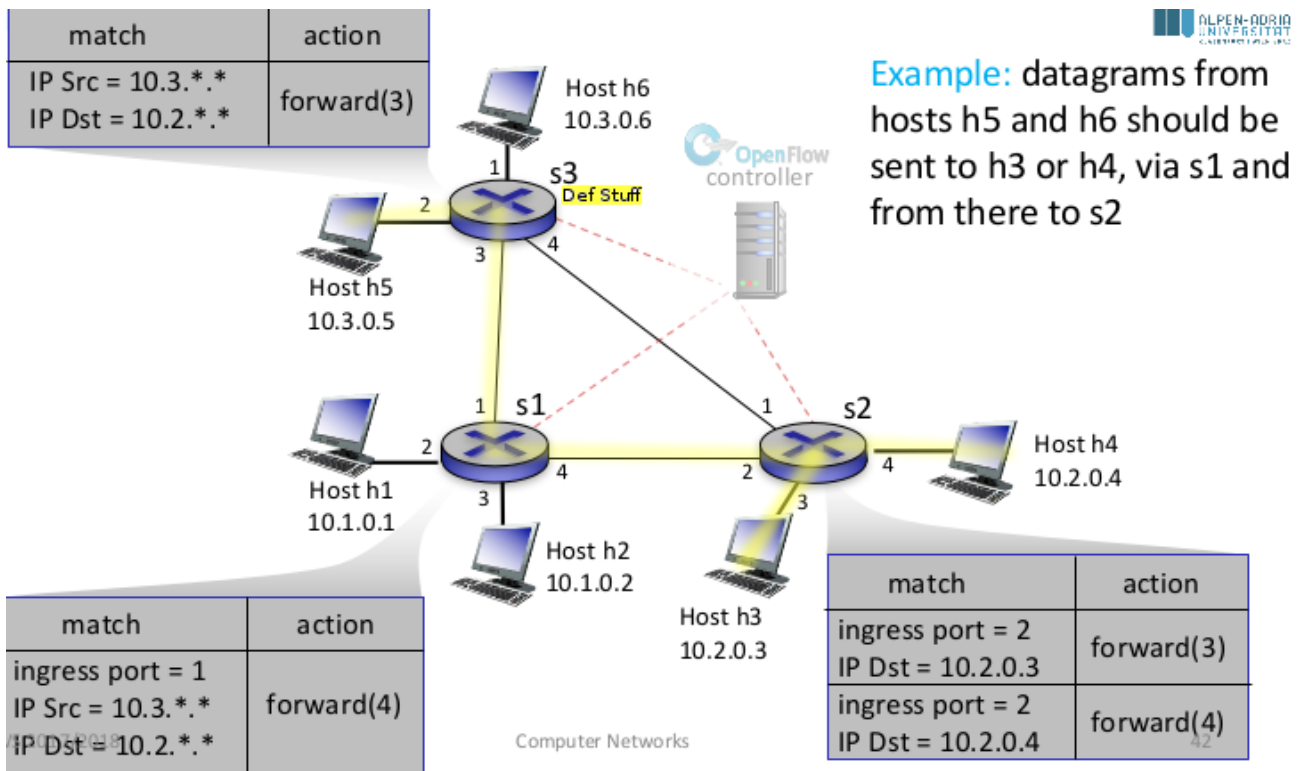
*layer 2 frames from MAC address 22:A7:23:11:E1:02 should be forwarded to output port 3*

definierbar in der OpenFlow Abstraction:

## OpenFlow Abstraction

**Match + action: unifies different kinds of devices**

- **Router**
  - Match: **longest destination IP prefix**
  - Action: **forward** to a link
- **Switch**
  - Match: **destination MAC address**
  - Action: **forward or flood**
- **Firewall**
  - Match: **IP addresses and TCP/UDP port numbers**
  - Action: **permit or deny**
- **NAT**
  - Match: **IP address and port**
  - Action: **rewrite address and port**



## Network Layer – Control Plane

# Network Layer Functions

Recall: **two network-layer functions**

- **Forwarding**: move packets from router's input to appropriate router output data plane
- **Routing**: determine route taken by packets from source to destination control plane
- **Two approaches** to structuring **network control plane**:
  - Per-router control (traditional)
  - Logically centralized control (software defined networking)


Forwarding: Packet kommt → Wohin gehts? Routing: Packet: Definiere Wohin. (Flow table)

# Per-Router Control Plane

Individual routing algorithm components in each and every router interact in the control plane

Values in arriving packet header

Definiert per Header wohin output geht.  
Routing Algs Arbeiten mit Flow Table



WS 2017/2018

Computer Net

4




## Logically Centralized Control Plane



A distinct (typically remote) controller interacts with local control agents (CAs)

Values in arriving packet header



WS 2017/2018

Com

5

- Data Flow via Control Agents definiert via Remote Controller.
- Kann physisch verteilt sein
- Einfach implementierbar.



## Routing Protocols

**Routing protocol goal:** determine “good” paths (equivalently, routes), from sending hosts to receiving host, through network of routers

- **Path:** sequence of routers packets will traverse in going from given initial source host to given final destination host
- **“Good”:** least “cost”, “fastest”, “least congested”
- Routing: a “top-10” networking challenge!

## Graph Abstraction of the Network

- graph:  $G = (N, E)$



$c(x, x') = \text{cost of link } (x, x')$   
e.g.,  $c(w, z) = 5$

Cost could always be 1, or inversely related to bandwidth, or inversely related to congestion

Cost of path  $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

**Key question:** what is the least-cost path between u and z?

**Routing algorithm:** algorithm that finds that least cost path



# Routing Algorithm Classification

---

## Q: global or decentralized info?

### Global:

- All routers have **complete topology**, link cost info
- "**Link state**" algorithms

### Decentralized:

- Router knows **physically-connected neighbors**, link costs to neighbors
- **Iterative process** of computation, exchange of info with neighbors
- "**Distance vector**" algorithms

### Local:

- Router **knows only itself**
- **Hot potato, flooding**, ...

## Q: static or dynamic?

### Static:

- **Routes change slowly over time**

### Dynamic:

- **Routes change more quickly**
  - Periodic update
  - In response to link cost changes

Wie: Protokollimplementierung

- Decentral: Distance Vector Algorithm (Know your neighbour); Link State Algorithm(Dijkstra)
- Lokal: Hot potatoe, Flooding (Lenkt an alle weiter == Broadcasting im THCP)
- Static: Linkweiterleitungen ändern sich nicht
- Dynamisch: Linkweiterleitungen ändern sich oft ↔ Kostenneuberechnung



# A Link-State Routing Algorithm

## Dijkstra's algorithm

- Network topology, link costs **known to all nodes**
  - Accomplished via “**link state broadcast**”
  - All nodes have **same info**
- Computes **least cost paths** from one node (source) to all other nodes
  - Gives **forwarding table** for that node
- **Iterative**
  - After **k iterations**, know least cost path to k destinations

Alle Knoten kennen ihre Flows.


## Dijkstra's Algorithm: Example

| Step | N'     | D(v)<br>p(v) | D(w)<br>p(w) | D(x)<br>p(x) | D(y)<br>p(y) | D(z)<br>p(z) |
|------|--------|--------------|--------------|--------------|--------------|--------------|
| 0    | u      | 7,u          | <b>3,u</b>   | 5,u          | ∞            | ∞            |
| 1    | uw     | 6,w          |              | <b>5,u</b>   | 11,w         | ∞            |
| 2    | uwx    | <b>6,w</b>   |              |              | 11,w         | 14,x         |
| 3    | uwxv   |              |              | <b>10,v</b>  |              | 14,x         |
| 4    | uwxvy  |              |              |              | <b>12,y</b>  |              |
| 5    | uwxvyz |              |              |              |              |              |

### Notes:

Construct shortest path tree by tracing predecessor nodes

Ties can exist (can be broken arbitrarily)



WS 2017/2018

Computer Networks

12

Bei Gleichheit : Wurscht, beide Varianten möglich; Dennoch definiert Router:

- Gibts zusätzliche Kriterien: #Hops, Zufall, Alphabetisch, ...

## Dijkstra's Algorithm

### Notation

$c(x,y)$ : link cost from node  $x$  to  $y$ ;  $= \infty$   
if not direct neighbors

$D(v)$ : current value of cost of path from  
source to destination  $v$

$p(v)$ : predecessor node along path  
from source to  $v$

$N'$ : set of nodes whose least cost path  
definitively known

### 1 Initialization:

2  $N' = \{u\}$

3 for all nodes  $v$

4 if  $v$  adjacent to  $u$

5 then  $D(v) = c(u,v)$

6 else  $D(v) = \infty$

7

### 8 Loop

9 find  $w$  not in  $N'$  such that  $D(w)$  is a minimum

10 add  $w$  to  $N'$

11 update  $D(v)$  for all  $v$  adjacent to  $w$  and not in  $N'$   
:

12  **$D(v) = \min( D(v), D(w) + c(w,v) )$**

13 /\* new cost to  $v$  is either old cost to  $v$  or known


14 shortest path cost to  $w$  plus cost from  $w$  to  $v$  \*/

15 **until all nodes in  $N'$**

# Dijkstra's Algorithm: Another Example

| Step | N'     | D(v),p(v) | D(w),p(w) | D(x),p(x) | D(y),p(y) | D(z),p(z) |
|------|--------|-----------|-----------|-----------|-----------|-----------|
| 0    | u      | 2,u       | 5,u       | 1,u       | ∞         | ∞         |
| 1    | ux     | 2,u       | 4,x       |           | 2,x       | ∞         |
| 2    | uxy    | 2,u       | 3,y       |           |           | 4,y       |
| 3    | uxyv   |           | 3,y       |           |           | 4,y       |
| 4    | uxyvw  |           |           |           |           | 4,y       |
| 5    | uxyvwz |           |           |           |           |           |


\* Check out the online interactive exercises for more examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)



# Dijkstra's Algorithm: Example (2)

Resulting shortest-path tree from u:

Resulting forwarding table in u:




| destination | link  |
|-------------|-------|
| v           | (u,v) |
| x           | (u,x) |
| y           | (u,x) |
| w           | (u,x) |
| z           | (u,x) |

→ Mögl. Wie man an Forwarding-Table kommt.

**Komplexität:**

# Dijkstra's Algorithm: Discussion

- Algorithm **complexity**:  $n$  nodes
  - Each iteration: **need to check all nodes**,  $w$ , not in  $N$
  - $n(n+1)/2$  comparisons:  $O(n^2)$
  - More **efficient implementations** possible:  $O(n \log n)$
- Oscillations** possible:
  - E.g., support link cost equals amount of carried traffic:



Oszillationen möglich :: Wenn Kostenmenge = Datenmenge:: 1 für D, e für C, 1 für B  
 → Loops. Möglichkeit: Keine Abhängigkeit von Datenvolumen ↔ Widerspruch zur Annahme  
 → Dezentraler Algorithmusverlaufs – Jeder Router für sich.

## Distance Vector Algorithm


### Bellman-Ford equation (dynamic programming)

let

$d_x(y) := \text{cost of least-cost path from } x \text{ to } y$


then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$



- “Know your neighbour” - Bellman-Ford equation
- Suche Minimale Kosten zu allen Nachbarn + Kosten vom Nachbarn zur Dest.

# Distance Vector Algorithm: Example



Clearly,  $d_v(z) = 5$ ,  $d_x(z) = 3$ ,  $d_w(z) = 3$

Bellman-Ford equation says:

$$d_u(z) = \min \{ c(u,v) + d_v(z), \\ c(u,x) + d_x(z), \\ c(u,w) + d_w(z) \} \\ = \min \{ 2 + 5, \\ 1 + 3, \\ 5 + 3 \} = 4$$

Node achieving minimum is next hop in shortest path, used in forwarding table

$$d_v(z) = 2 + 1 + 2$$

$$d_x(z) = 1 + 2$$

$$d_w(z) = 1 + 2$$

$$c(u,v) = 2$$

$$c(u,x) = 1$$

KLINGENSCHULTZ 1 W. ED. 6902

# Distance Vector Algorithm

- $D_x(y)$  = estimate of least cost from x to y
  - x maintains distance vector  $D_x = [D_x(y) : y \in N]$
- Node x:
  - Knows cost to each neighbor v:  $c(x,v)$
  - Maintains its neighbors' distance vectors. For each neighbor v, x maintains  $D_v = [D_v(y) : y \in N]$

## Key idea:

- From time-to-time, each node sends its own distance vector estimate to neighbors
- When x receives new DV estimate from neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{ c(x,v) + D_v(y) \} \text{ for each node } y \in N$$

- Under minor, natural conditions, the estimate  $D_x(y)$  converge to the actual least cost  $d_x(y)$

... Maintain routes logical costs. Update periodically

Iterative, asynchronous: each local iteration caused by:

- Local **link cost change**
- DV **update message** from neighbor

Distributed:

- Each node **notifies neighbors only** when its DV changes
  - Neighbors then notify their neighbors if necessary

## • Each node


**Wait** for (change in local link cost or msg from neighbor)

**Recompute** estimates

If DV to any destination has changed, **notify** neighbors



→ Bei Änderung innerhalb eines Knotens, so benachrichtige alle Nachbarn → Welche wiederum Nachbarn benachrichtigen.

### ► BITMOVIN



$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\} \\ = \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\} \\ = \min\{2+1, 7+0\} = 3$$



018


Computer Networks

21


## Distance Vector: Link Cost Changes

### Link cost changes

- Node detects **local link cost change**
- Updates routing info, **recalculates distance vector**
- If DV changes, **notify neighbors**





"Good news travel fast"





## Link cost changes



- Node detects **local link cost change**
- Bad news travels slow – "**count to infinity**" problem!
- 44 iterations before algorithm stabilizes



^ Umgehbar mit Poisoned Reverse:

## Poisoned reverse

- If **Z routes through Y** to get to X:
- Z tells Y its (Z's) distance to X is **infinite** (so Y won't route to X via Z)
- Will this completely solve count to infinity problem?



# Comparison of LS and DV algorithms

---

## Message complexity

- LS: with  $n$  nodes,  $E$  links,  $O(nE)$  messages sent
- DV: exchange between neighbors only
  - Convergence time varies

## Speed of convergence

- LS:  $O(n^2)$  algorithm requires  $O(nE)$  messages
  - May have oscillations
- DV: convergence time varies
  - May be routing loops
  - Count-to-infinity problem

## Robustness: what happens if router malfunctions?

- LS:
  - Node can advertise incorrect link cost
  - Each node computes only its own table
- DV:
  - DV node can advertise incorrect path cost
  - Each node's table used by others
  - Error propagate thru network

LS: Kennt alle Kosten (?)

DV: Kennt nur Nachbar

→ Beide Arten sind 'idealisiert'

# Making Routing Scalable

---

- Our routing study thus far - idealized
  - all routers identical & network "flat"... not true in practice

- Scale: with billions of destinations:
  - Can't store all destinations in routing tables!
  - Routing table exchange would swamp links!

- Administrative autonomy

- Internet = network of networks

- Each network admin may want to control routing in its own network

Aggregate routers into regions known as "autonomous systems" (AS) (a.k.a. "domains")

## Intra-AS Routing

- Routing among hosts, routers in same AS ("network")
- All routers in AS must run same intra-domain protocol
- Routers in different AS can run different intra-domain routing protocol
- Gateway router: at "edge" of its own AS, has link(s) to router(s) in other AS'es

## Inter-AS Routing

- Routing among AS'es
- Gateways perform inter-domain routing (as well as intra-domain routing)

